

# Real time quiz feature engineering proposal

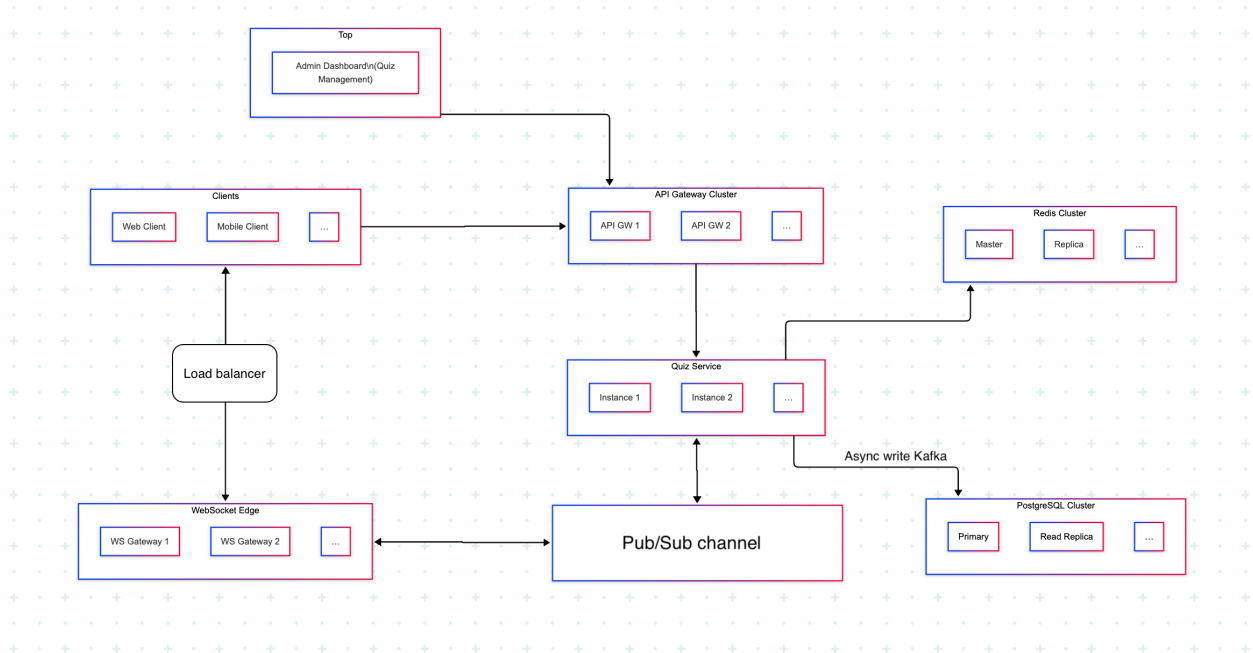
## Context

Create a technical solution for a real-time quiz feature for an English learning application. This feature will allow users to answer questions in real-time, compete with others, and see their scores updated live on a leaderboard.

## Menu

- **Architecture Overview**
- **Data Flow**
- **Technology Justification**
- **Engineering Specifications**
- **Trade-offs Analysis**
- **Observability & Monitoring**
- **Client Architecture**
- **Rollout Strategy**
- **Rollback Plan**
- **Testing**
- **Readiness Checklist**

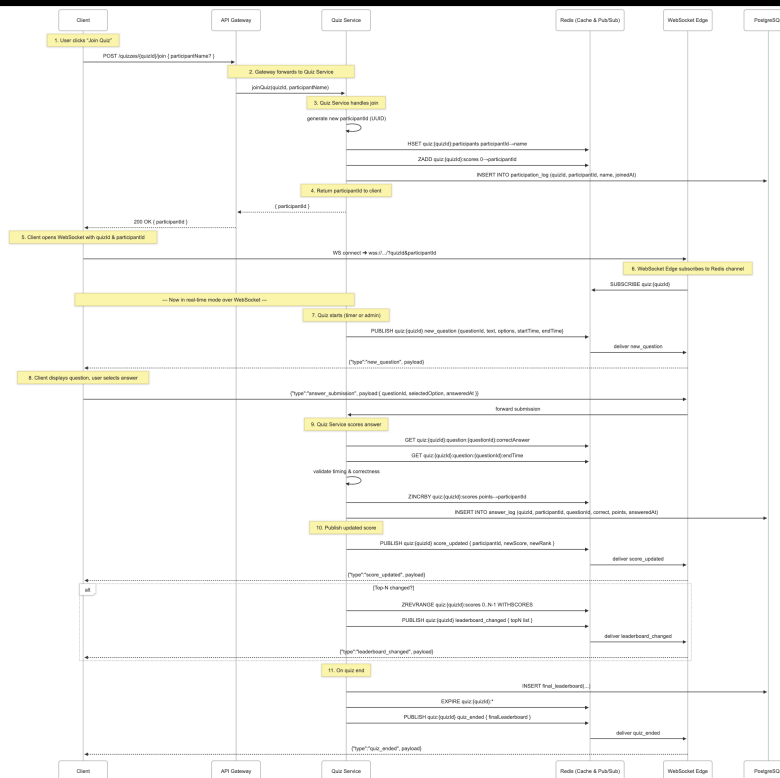
## Diagram



## Component Descriptions

Component	Role
<b>Admin Dashboard</b>	UI for creating quizzes, defining questions, and monitoring active sessions.
<b>API Gateway Cluster</b>	Exposes REST endpoints (join, metadata), handles auth, rate-limits.
<b>Quiz Service</b>	Core logic: validates joins, sequences questions, scores answers, and publishes events. Stateless & horizontally scalable.
<b>Redis Cluster</b>	In-memory store for active session state (participants, scores, current question) and Pub/Sub broker for real-time events.
<b>WebSocket Edge Cluster</b>	Maintains client socket connections, subscribes to Redis channels, and pushes real-time messages (questions, scores, leaderboard) to participants.
<b>Clients</b>	Web or mobile apps that (a) REST-join the quiz, (b) open one WebSocket, (c) render questions & leaderboards, (d) send answer messages.
<b>PostgreSQL Cluster</b>	Durable storage of quiz definitions, historical logs (joins, answers, final scores) for audit and analytics.

# Data flow



## • User Joins the Quiz

- The participant enters a quiz ID and clicks "Join."
- A HTTP request lands at API gateway, which creates a unique `participantId` for them, stores it in our fast in-memory store (Redis), and records the join in our database.

## • WebSocket Connection Established

- Armed with their new `participantId`, the client immediately opens a single persistent WebSocket connection, identifying itself by quiz and participant IDs.
- Behind the scenes, that socket is tied to a Redis Pub/Sub channel for this quiz.

- **Question Broadcast**

- When the quiz is ready to start when participants > 2 the quiz will start in 5 seconds (my product design), the core Quiz Service publishes a "new question" event into Redis.
- Our WebSocket layer picks that up and instantly pushes the question (text, options, start/end times) to every connected participant.

- **Answer Submission Over Socket**

- Each user selects their answer and sends an "answer submission" message over the same WebSocket.
- That message carries the `questionId`, chosen option, and a timestamp.

- **Scoring & Immediate Feedback**

- The Quiz Service validates the answer (correctness and timing), then updates the user's score atomically in Redis.
- It also logs the detail in PostgreSQL for auditing.

- **Real-Time Score Update**

- As soon as the score changes, the WebSocket layer relays this in under 100 ms to **all** participants, so everyone sees the updated score and rank.
- If the top standings shift, the service pulls the top-N list from Redis and publishes a "leaderboard changed" event.
- Again, WebSocket instantly rebroadcasts that to every client.

- **Quiz Completion & Cleanup**

- At the end, final scores are written to our durable store, Redis session data expires, and a final "quiz ended" event delivers the conclusive leaderboard to all participants.

## Technology Justification

Concern	Chosen Tool(s)	Why This Choice?
<b>REST API</b>	API Gateway (e.g. AWS/GCP)	Centralized auth, throttling, and routing; familiar HTTP patterns for

		simple join operation.
<b>Business Logic</b>	Quiz Service (Spring Boot)	Stateless microservice; easy horizontal scaling; rich ecosystem for Redis & Postgres integration.
<b>In-Memory State &amp; Pub/Sub</b>	Redis Cluster (with Pub/Sub)	Ultra-low latency for score updates; sorted sets for real-time leaderboards; built-in pub/sub.
<b>Real-Time Delivery</b>	WebSocket Edge (Socket.IO)	Single persistent connection per client; scales horizontally; decoupled from quiz logic via Redis.
<b>Persistent Storage</b>	PostgreSQL (Primary + Replicas)	ACID guarantees for quiz definitions and audit logs; read replicas support dashboard/reporting.
<b>Deployment &amp; Scaling</b>	Kubernetes / Managed Containers	Automated scaling, rolling updates, multi-AZ deployments for reliability and maintainability.
<b>Monitoring &amp; Logging</b>	Prometheus + Grafana; ELK / Sentry	Metrics and dashboards for health & performance; centralized logs and error tracking.

## Engineering Specification

### 1. POST /api/v1/quizzes/{quizId}/join

- **Description:** Register a participant for a quiz session.
- **Path Parameter:**
  - `quizId` (string, required): Unique quiz identifier.
- **Request Body** (application/json):

```
{
  "participantName": "string"
}
```

- **Response** (200 OK, application/json):

```
{
  "quizId": "string",    // echoed
  "participantId": "string" // server-generated UUID
}
```

## WebSocket Contract

Clients open a single WebSocket to: `wss://<host>/ws?quizId={quizId}&participantId={participantId}`

All messages are JSON with fields:

```
{
  "type": "string",    // event type
  "payload": { ... }   // event-specific data
}
```

### Client → Server Events

#### 1. **answer\_submission**

- **When:** User submits an answer.
- **Payload:**

```
{
  "questionId": "string",
  "selectedOption": "string",
  "answeredAt": "ISO8601 timestamp"
}
```

#### 2. **get\_leaderboard** (optional)

- **When:** Client needs a full snapshot over WS.
- **Payload:** `{}`

#### 3. **heartbeat** (optional)

- **When:** Keep-alive ping.

- **Payload:** `{}`

## Server → Client Events

### 1. new\_question

- **When:** Quiz starts or moves to next question.
- **Payload:**

```
{
  "questionId": "string",
  "text": "string",
  "options": ["string", ...],
  "startTime": "ISO8601 timestamp",
  "endTime": "ISO8601 timestamp"
}
```

### 2. leaderboard\_changed

- **When:** Top-N standings change.
- **Payload:**

```
{
  "leaderboard": [
    {"participantId": "string", "score": number, "rank": integer}
  ]
}
```

### 3. leaderboard\_data

- **When:** Response to `get_leaderboard`.
- **Payload:** Same structure as `leaderboard_changed`.

### 4. quiz\_ended

- **When:** Quiz concludes.
- **Payload:**

```
{
  "finalLeaderboard": [
    {"participantId":"string","score":number,"rank":integer}
  ]
}
```

## 5. error

- **When:** Validation or server error.
- **Payload:**

```
{
  "reason": "string",
  "details": { ... }
}
```

## JSON Schema References

- **Participant Object:**

```
{ "type": "object",
  "properties": { "participantId": { "type": "string" }, "participantName": { "type": "string" }, "score": { "type": "number" }, "rank": { "type": "integer" } }, "required": [ "participantId", "score", "rank" ] }
```

- **Leaderboard Array:** `{ "type": "array", "items": { "$ref": "#/definitions/Participant" } }`

## Trade-offs

### 1. Redis Pub/Sub vs. Durable Messaging (e.g. Kafka)

Trade-off	Description
<b>Pros</b>	Redis Pub/Sub is extremely fast and simple to integrate for fanout (WS delivery). Low latency (<5ms).



<b>Cons</b>	No delivery guarantees. No message history or replay. If a WS node is down, it <b>misses the message</b> .
<b>Why Acceptable</b>	WS delivery is real-time and ephemeral—missing 1 leaderboard update is not critical. Simplicity and speed were prioritized.

## 2. In-memory leaderboard with Redis vs. recomputing from DB

Trade-off	Description
<b>Pros</b>	Redis sorted sets ( <b>ZSET</b> ) make score updates and leaderboard queries extremely fast ( $O(\log N)$ ).
<b>Cons</b>	Redis memory can become a bottleneck (especially on hot keys). Must manually enforce TTL or cleanup.
<b>Why Acceptable</b>	Redis is highly optimized for this pattern. We accept higher memory usage to keep leaderboard performance real-time at scale.

## 3. Async DB writes (Kafka) vs. immediate DB consistency

Trade-off	Description
<b>Pros</b>	Offloads write pressure from <b>QuizService</b> . Enables high-QPS submissions without overloading PostgreSQL.
<b>Cons</b>	DB is eventually consistent. Data isn't written immediately after submission. Cannot <b>SELECT</b> answers synchronously.
<b>Why Acceptable</b>	UI and real-time experience do not depend on DB writes. Consistency delay of seconds is fine for post-game analysis.

## 4. Single leaderboard (top-N) vs. per-participant ranking requests

Trade-off	Description
<b>Pros</b>	Top-N is cheap to cache and broadcast to all clients.
<b>Cons</b>	If you want to show "You're ranked #847" to each user, you'd need per-user <b>ZRANK</b> , which is expensive at scale.
<b>Why Acceptable</b>	Tradeoff was made in favor of shared leaderboard UX. Per-user rank can be fetched less frequently or delayed.

## 5. WebSocket delivery only vs. fallback polling

Trade-off	Description
<b>Pros</b>	Keeps client UX real-time and reactive. Low latency push via WS.
<b>Cons</b>	If WS fails or lags, client has no fallback mechanism to poll for state.
<b>Why Acceptable</b>	WS is monitored, and quiz state is ephemeral. For resilience, fallback polling <b>can</b> be added later.

## Summary Table

Area	Tradeoff Made	Chosen for
Messaging	Redis Pub/Sub vs Kafka	Simplicity, low latency
Leaderboard	Redis vs DB recompute	Real-time responsiveness
Writes	Async via Kafka vs direct DB	Throughput scalability
Ranking model	Top-N vs Per-user rank	Fanout efficiency
Client delivery	WebSocket only vs fallback	Real-time UX

## Observability & Monitoring

To ensure the quiz platform operates reliably at scale (up to 1M concurrent users), we apply full-stack observability across services, infrastructure, and real-time flows.

This enables:

- Fast incident detection and diagnosis
- Performance tuning and capacity planning
- Auditability of key events (joins, answers, broadcasts)

## Observability Stack

Pillar	Tools
Logs	Loki, ELK (Elasticsearch + Logstash + Kibana), or GCP Cloud Logging
Metrics	Prometheus + Grafana

Tracing	OpenTelemetry + Jaeger or Zipkin
Alerts	Prometheus Alertmanager → Slack, PagerDuty

All services expose `/actuator/prometheus` or native exporters.

## Instrumented Components

### Quiz Service

- `@Timed` on `joinQuiz()`, `submitAnswer()`
- Counter: `answers_submitted_total{quizId}`
- Timer: `quiz_submission_latency_seconds`
- Kafka producer errors
- Leaderboard update timing

### WebSocket Edge

- Active WS connections
- `ws_outgoing_bytes_total`
- `ws_connection_errors_total`
- Redis pub/sub latency (if applicable)

### Redis

- `redis_connected_clients`
- `redis_used_memory_bytes`
- `redis_pubsub_channels`
- `redis_commands_duration_seconds`

### Kafka

- `kafka_produce_errors_total`
- `kafka_consumergroup_lag`
- `kafka_batch_flush_latency_seconds`

## PostgreSQL

- Write QPS
- Replication lag
- Connection pool usage

## Dashboards

Category	Panel Example
WebSocket Edge	Active connections per node
Quiz Events	Join rate, submission rate, avg answer latency
Redis	Command rate, memory, slowlogs
Kafka	Per-topic lag, consumer group throughput
PostgreSQL	Insert throughput, slow queries, replication
Errors	4xx/5xx breakdown, Redis failures

## Alerting Rules (via Alertmanager)

Alert Condition	Description
<code>redis_used_memory &gt; 80%</code>	Redis capacity pressure
<code>kafka_consumergroup_lag &gt; 10s for 1m</code>	Async writes falling behind
<code>quiz_submission_latency_seconds &gt; 500ms P95</code>	QuizService under load
<code>ws_connection_errors_total</code> spike	Broken WS layer or Redis link
<code>http_5xx_total &gt; 1% over 5m</code>	Service instability

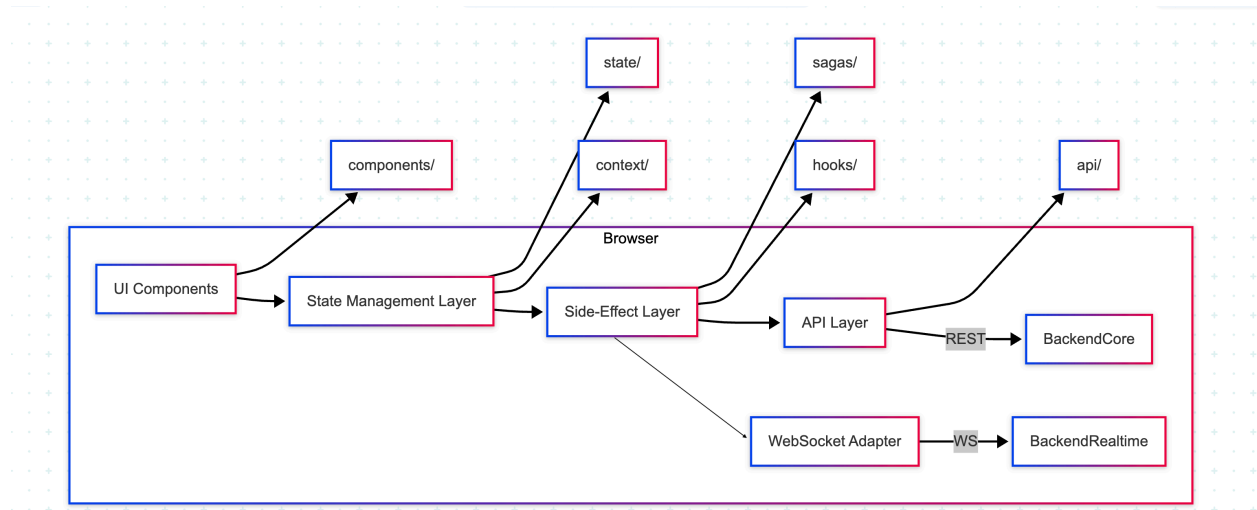
Alerts route to Slack or PagerDuty based on severity.

## Logging Strategy

- Use **JSON structured logs**
- Add `quizId`, `participantId`, and `questionId` as contextual fields (careful with cardinality)
- All services forward logs to Loki or ELK

- WS and Quiz logs correlated via request IDs or socket session IDs

## Client Architecture



## Component Descriptions

Layer / Folder	Description
<b>components/</b>	Presentation layer. Includes <code>QuizPage</code> , <code>Leaderboard</code> , <code>QuestionCard</code> , etc. Renders state and emits user interactions (e.g., join, submit answer).
<b>state/</b>	Redux Toolkit slice ( <code>quizSlice</code> ) defines the global state: quiz ID, participant info, current question, leaderboard, and errors.
<b>sagas/</b>	Implements real-time event handling: joins quiz, opens WebSocket, listens to <code>new_question</code> and <code>leaderboard_changed</code> , and handles answer submission.
<b>api/</b>	Wraps REST API calls such as <code>joinQuiz(quizId, name)</code> , abstracting fetch logic from the UI.
<b>BackendCore</b>	Exposes REST endpoints ( <code>/join</code> ) to register participants and submit answers.

## BackendRealtime

WebSocket server that pushes messages like `new_question`, `leaderboard_changed`, and receives events like `answer_submission`.

# Data Flow: User Joins → Answers → Leaderboard Updates

## A. User Joins a Quiz

### 1. UI Action:

- User submits name and quiz ID via `QuizPage`.

### 2. Redux Dispatch:

- `joinRequest({ quizId, name })` is dispatched.

### 3. Saga Effect:

- `handleJoin()` saga calls `joinQuiz()` API (REST).
- On success, dispatches `joinSuccess({ quizId, participantId, name })`.

### 4. WebSocket Connection:

- Saga opens a WebSocket:

```
ws://localhost:8090/ws/quiz?quizId=...&participantId=...
```

- Sets up `eventChannel(ws)` to listen for messages.
- Forks a listener loop:
  - If `msg.type === 'new_question'` → dispatch `receiveQuestion()`
  - If `msg.type === 'leaderboard_changed'` or `leaderboard_data` → dispatch `receiveLeaderboard()`

## B. New Question Broadcast

1. **BackendRealtime** pushes `new_question` via WebSocket.
2. **WebSocket Adapter** receives event.

3. **Saga** catches it → dispatches `receiveQuestion(payload)` .
  4. **Redux State** is updated → UI ( `QuestionCard` ) re-renders.
- 

## C. User Submits Answer

### 1. UI Action:

- User selects an option → dispatches `submitAnswer({ questionId, selectedOption })` .

### 2. Saga Effect:

- `handleSubmitAnswer()` selects `participantId` from Redux.
- Constructs message:

```
{
  "type": "answer_submission",
  "payload": {
    "participantId": "...",
    "questionId": "...",
    "selectedOption": "...",
    "answeredAt": "..."
  }
}
```

- Sends via WebSocket.
- 

## D. Leaderboard Update

1. **BackendRealtime** emits `leaderboard_changed` via WebSocket (after scoring logic).
2. **Saga WebSocket listener**:
  - Catches the message → dispatches `receiveLeaderboard()` .
3. **Redux State** updates `leaderboard` array.
4. **Leaderboard UI** re-renders instantly.

## Reliability – Retry Strategy

To ensure robustness and seamless user experience during network disruptions or transient backend failures, we implement a **retry strategy** for both the initial quiz join request ( `joinQuiz` API call) and the WebSocket connection. This helps minimize user-facing errors during periods of instability while providing visual feedback about ongoing reconnection attempts.

---

## Goals

- Handle temporary failures (e.g., network hiccups, backend downtime).
  - Provide automatic retries with exponential backoff.
  - Visibly inform users of retry progress and eventual success or failure.
  - Avoid infinite loops or aggressive reconnect attempts.
- 

## Retry Logic for `joinQuiz` API

The `joinQuiz(quizId, name)` call is wrapped with Redux-Saga's built-in `retry` effect:

```
const res: { quizId: string; participantId: string } = yield retry(3, 1000, joinQuiz, quizId, name);
```

- **Retries:** 3
- **Interval:** 1000 ms between attempts
- **Backoff:** Linear
- **Failure Mode:** Falls through to `joinFailure(err.message)` if all retries fail

This ensures a reasonable number of attempts before presenting an error to the user.

---

## Retry Logic for WebSocket Connection

WebSocket reconnection is handled manually using an exponential backoff strategy. The logic attempts to connect up to 5 times with increasing delays:



```
function* connectWebSocketWithRetry(url: string, maxRetries = 5)
```

- **Retries:** 5
- **Backoff:**  $2^{\text{attempts}} * 100$  ms (200ms, 400ms, 800ms, etc.)
- **Failure Mode:** Throws error after final attempt

The implementation ensures that the socket is successfully opened before continuing. On failure, the user is notified through the UI.

---

## Resilience & Test Strategy

To ensure production-grade reliability, we adopt a layered test strategy:

### Automated Tests

- Unit tests for QuizService logic (joining, scoring, ranking)
- Integration tests for Redis, Kafka, PostgreSQL flows
- Contract tests for WebSocket message schema (consumer/producer validation)

### Real-Time Simulation

- Simulate 10k clients submitting answers via [locust.io](https://locust.io) or k6 with WS support
- Verify leaderboard accuracy under load using fuzzed answer inputs

### Chaos & Fault Injection

- Redis latency injection using [toxiproxy](https://toxiproxy.net)
- Kafka consumer kill-switch scenarios
- WS disconnection / reconnection tests
- Region-specific outages simulated via Kubernetes failure injections

## Validation Goals

- Score consistency
- Leaderboard ranking under churn
- Graceful degradation with missing WS events

## Rollout Plan

To safely and incrementally roll out the real-time quiz platform to production with confidence. This plan minimizes user impact, validates system behavior under load, and ensures rollback is quick if needed.

---

### Deployment Phases

#### Phase 0 – Internal QA

Scope	Details
Environment	Internal staging (simulated load)
Users	Internal team only
Scale	100–500 concurrent clients
Monitoring	Full observability stack validated
What to Validate	Leaderboard sync, question timing, scoring

---

#### Phase 1 – Beta Group (Shadow Traffic)

Scope	Details
Environment	Production infrastructure
Users	Opt-in beta testers or staff
Scale	1–5k concurrent
What to Validate	Real traffic pattern, Redis scaling, WS fanout
Technique	Feature flag / allowlist
Safety	Logs, alerts, and metrics under scrutiny

---

## Phase 2 – Gradual Rollout

Step	% of Users	Actions
Step 1	10%	Monitor for WS stability, Kafka lag
Step 2	25%	Observe Redis memory usage, DB write spikes
Step 3	50%	Validate leaderboard sync across regions
Step 4	100%	Full rollout with on-call active

Each step is gated by:

- No SLO violations
- Error rate < 0.5%
- Kafka lag < 5s
- Redis memory < 70%

## Multi-Region Deployment

To support global audiences with low-latency access:

### Strategy

- **Redis Cluster:** Deployed in each region as active-active clusters. Quiz sessions are region-localized via quiz ID prefixing (e.g., `eu-quiz123`).
- **WebSocket Edge:** Deployed per region using Anycast DNS or CDN-based routing (e.g., Cloudflare Spectrum, AWS Global Accelerator).
- **API Gateway:** Multi-region, with global routing based on latency (e.g., GCP Load Balancer or AWS ALB with geo-routing).

### Quiz Session Affinity

Users join region-local quiz sessions. Cross-region play is not supported in MVP but can be enabled with a Global Redis cache layer or session replication.

### Future Work

- “Quiz Session Handoff” between regions (for future cross-region support)

- “Distributed Consistency Testing” for Redis replication correctness

## Latency Budget

To maintain a responsive real-time experience, we define the following latency budgets per operation (P95 targets):

Operation	Latency Target (P95)
REST Join API	< 150 ms
WebSocket connection setup	< 200 ms (TLS + Auth)
Submit Answer → Scoring	< 100 ms (end-to-end)
Score → Leaderboard Update	< 80 ms (PubSub + WS fanout)
Quiz End → Final Leaderboard push	< 200 ms

All metrics are monitored via Prometheus histograms and visualized on Grafana dashboards. Alerts are configured to detect spikes beyond budget.

## Rollback Plan

Trigger Condition	Rollback Action
Redis saturation (>90%)	Scale Redis cluster / purge keys
Kafka consumer lag > 30s	Pause traffic → investigate
PostgreSQL slow inserts / timeouts	Fallback to write-buffer only mode
WS nodes crashloop / mem spike	Reduce connection cap per node
P99 latency > 500ms across endpoints	Roll back to previous deployment

Rollback is immediate via:

- Blue/green or canary deploy with toggles
- Helm rollback or CI/CD revert
- Feature flag switch-off

## Readiness Checklist

Area	Ready?	Notes
------	--------	-------

Redis scaling	✓	Clustered, key TTL, hot quiz sharding
Kafka buffering	✓	Write path async, consumers monitored
DB write safety	✓	Batched or buffered via stream
WS connection	✓	Load balanced, fanout from pub/sub
Logging	✓	JSON structured, central aggregation
Metrics	✓	Prometheus + Grafana alerts configured
Alerting	✓	PagerDuty / Slack routing in place